

Zcash FPGA acceleration engine

Version 1.4 release

Ben Devlin

bsdevlin@gmail.com

GitHub repo: <https://github.com/bsdevlin/zcash-fpga/>

Version history	3
Terms used	4
Overview	5
Zcash FPGA project	5
Interfaces and FPGA hardware	5
Project goals	6
Phase 1	6
Phase 2	6
Implementation	7
Overview	7
FPGA Memory Map	7
Streaming commands	7
SW to FPGA	8
reset_fpga	8
get_fpga_status	8
verify_equihash	8
verify_secp256k1_sig	8
FPGA to SW	9
reset_fpga_rpl	9
fpga_status_rpl	9
fpga_ignore_rpl	10
verify_equihash_rpl	10
verify_secp256k1_sig_rpl	10
bls12_381_interrupt_rpl	10
FPGA command capability register	11
FPGA Architecture	11
Overview of blocks in the system	11
Interface module	12
AWS (Amazon)	12
VHH (Bittware)	12
Equihash Verification Engine	12
Overview	12
Block diagram	13

Performance evaluation	13
FPGA resources	13
Clock cycles	13
Transparent Signature Verification Engine (secp256k1 ECDSA core)	14
Overview	14
Block diagram	15
Performance evaluation	15
FPGA resources	15
Clock cycles	15
Future Optimizations	16
BLS12-381 Coprocessor (zk-SNARK accelerator)	16
Overview	16
Instructions	17
Memory Map	18
Architecture	19
Performance Comparison	21
FPGA resources	21
Clock cycles	22
Future Optimizations	22
User Guide	23
Running Simulations	23
Module level simulations	23
AWS Board level	23
Usage with a local FPGA board	23
Usage on AWS	23
Building the FPGA image	24
Loading FPGA image	25
Rust interface	25
Startup test program	25
FPGA debug	27
Latest AFIs	27
Old (unused) AFIs	28
Conclusions	29
Appendix	31
Example decoding Zcash block #346	31

Version history

- v1.1
 - First major release of the code, includes many reusable logic cores, along with the equihash engine, secp256k1 signature verification engine, and bls12-381 coprocessor with Fp and Fp^2 point multiplication (pairing to be implemented in v1.2)
 - Top level module for the Zcash acceleration engine
 - Top level board files for both Bittware VVH and Amazon AWS EC2 F1 FPGAs
 - bls12-381 coprocessor so far has only been tested on AWS
 - Document still missing content for some sections, will be completed in v1.2
- v1.2
 - bls12-381 coprocessor
 - Added optimal ate pairing module (instruction ATE_PAIRING)
 - Removed some instructions that were not used for control
 - Updated diagrams and performance numbers
- v1.3
 - bls12-381 coprocessor
 - Added instructions for MILLER_LOOP and FINAL_EXP to allow for multi-pairing operations
 - Added Fp^{12} to the MULT instruction
 - Removed FPOINT_MULT instructions as they can all be covered by POINT_MULT
 - Changes to systemverilog to improve timing
 - Added AWS builds
 - Added sample output of test program
- v1.4
 - Added accum_mult_mod block to /ip_cores/ and modified bls12-381 coprocessor to use this
 - Modular multiplier using carry save adders for multiplication and RAM tables for reduction, 3X performance compared to previous Karatsuba multiplier

Terms used

FP (Field point)	FE (Field element)	JB (Jacobian)	AF (Affine)
FPGA (Field programmable gate array)	EC (Elliptic curve)	SW (Software - generally meaning what runs on the CPU)	AXI (Advanced eXtensible Interface)
Non-adjacent form (NAF)	RAM (Random Access memory)	BRAM (Block RAM) - on Xilinx FPGAs	URAM (Ultra RAM) - on Xilinx FPGAs
zk-SNARK (Zero-Knowledge Succinct Non-Interactive Argument of Knowledge)	ECDSA (Elliptic Curve Digital Signature Algorithm)	FFF (Fast Fourier transform)	

Overview

Zcash FPGA project

Zcash FPGA acceleration engine is a FPGA system used to accelerate the Zcash network. The **first phase** is focused on accelerating verification components of the blockchain, and the **second phase** is focused on zk-SNARK acceleration and elliptic curve operations required. All code developed is written in system verilog and open source under the GPL 3.0 license, intended to be modular and parameterizable for reuse, and can be found at the GitHub repo linked on the first page of this document.

FPGA acceleration allows us to offload work to a chip that is configured at the gate level to do specific hardware functions, and can bring several **advantages** over a CPU implementation:

1. Can be configured for large parallelism - e.g. you could configure an FPGA to do 1000x 32bit multiplications all at the same time allowing for large throughputs
2. Specialized functions that an x86 processor takes many instructions to implement could be implemented as a single instruction on an FPGA
3. Low latency direct access to data - e.g. you could develop custom TCP/IP hardware on an FPGA bypassing a NIC card / having a CPU make decisions

But also has **disadvantages**:

4. Clock speed is much slower on FPGA (100MHz - 300MHz depending on logic implemented) compared to a CPU (3GHz+) with multiple cores
5. Getting data in and out of the FPGA from the CPU takes roughly ~300ns(PCIe roundtrip) which translates to ~1000 clock cycles on a CPU even before we start processing
 - a. This is for an optimized core - AWS FPGAs used in this experiment take 1us+ roundtrip
6. Development cycle is much slower compared to CPU and not as easily accessible to a SW engineer

The goal of this work is not only to develop open source FPGA acceleration code for various Zcash systems and that can be of use to the wider community, but also to investigate/research the direction for future development (i.e. what cases are good candidates for acceleration and what cases are better left to SW).

Interfaces and FPGA hardware

The FPGA engine is designed to either be implemented on a Bittware board (VU37P FPGA w/ 8GB HBM, 16GB DDR4) or run on an Amazon AWS EC2 F1 FPGA instance (VU9P w/ 64GB DDR4). Both FPGAs are the same generation and speed grade, but depending on the board clock rates on FPGA might have to be scaled so that timing closure can be met (AWS FPGAs require extra “glue” logic and seem to not meet timing as easily as the VU37P). I have tried to use non-vendor specific blocks where possible (i.e. BRAMs, core logic, is mostly written from scratch in systemverilog), but in case cases I have used Xilinx IP for simplicity (mainly in the AWS top level, where the .xci files are included in the /ip folder). It would not take much work to implement the same code on an Altera FPGA or older generation Xilinx FPGA.

Communication to FPGA is split into two main methods:

1. Based on commands that are formatted with a header, followed by optional data (inputs for the command). FPGA sends replies to SW after a command is completed or in the case of any errors. These are sent over an AXI4-stream interface.

2. Using an instruction memory and data register approach, SW has direct access to FPGA memory and can configure more complex logic flows. Interrupt commands can be implemented so FPGA will send data to SW without required polling of FPGA memory. This is used for the bls12-381 coprocessor in phase 2. These are sent over an AXI4-lite interface.

Depending on the FPGA board used communication is either exposed to SW through a C++ library over PCIe (when using AWS), or over USB-UART (when using the Bittware board). There are wrappers that convert the communication method to the internal FPGA AXI-lite and AXI-stream interfaces.

Project goals

At a high level the FPGA architecture currently comprises several engines for dedicated tasks to handle the commands from SW, where more engines are to be added as development continues:

- Blake2b hash
- SHA256 hash
- Equihash verification engine
- Transparent signature verification engine (accelerate point multiplication on the secp256k1 curve)
- BLS12-381 coprocessor (accelerate EC operations on the bls12-381 curve such as point multiplication and pairing)

Phase 1

Phase 1 is focused on offloading various aspects of verifying the Zcash block chain onto the FPGA. These will include:

1. A equihash verification engine, which can take in a block header + solution and verify it is correct, as well as other fields in the block header that require processing (such as hashing)
2. Verifying transparent transaction in the block chain, which will be done by implementing a secp256k1 engine that can take in signatures and verify their correctness.

Phase 2

Phase 2 is focused on accelerating zk-SNARK operations.

1. This will be implemented a BLS12-381 coprocessor, where software can write instruction memory on the FPGA that will allow for chaining of multiple commands without having to send data in and out of the FPGA. This coprocessor will implement F_p , F_p^2 , F_p^6 , F_p^{12} arithmetic over the bls12-381 curve, as well as several higher level operations such as inversion, calculating powers, calculating frobenius map, miller loop, final exponentiation, and optimal ate pairing. Software can read and write both data and instruction memory to poll the current status of the coprocessor, or interrupt instructions can be used to send interrupts back to SW when certain commands complete.

The main goals for acceleration using this coprocessor*:

- Generate a shielded Zcash (Sapling) transaction with acceleration from the coprocessor
- Sign a shielded Zcash (Sapling) transaction with acceleration from the coprocessor

*The processor has been implemented on FPGA but requires some work in Zcash's Rust code to correctly interface with the FPGA and utilize, which has been made a lower priority at the moment for this project. A Rust wrapper around the cpp FPGA library has been developed, but still requires work before it could be released into production and used with an AWS F1 FPGA instance.

Implementation

Overview

- **FPGA:**
 - Bittware XUPV VH dev board w/ Virtex UltraScale+ VU37P HBM VCU128-ES1 (8GB HBM, 16GB DDR4)
 - Interface to host over UART (USB)
 - AWS EC2 F1 FPGA instance UltraScale+ VU9P (64GB DDR4)
 - Interface to host over PCIe
- **Software API:**
 - C++ library for AWS boards over PCIe - this is in the github repo `aws/cl_zcash/software/runtime/zcash_fpga.hpp`
 - A rust interface is in development and should be released in a later version, to allow the Zcash client to run on an AWS instance to utilize FPGA acceleration
 - USB-UART for Bittware boards using Python - this is in the github repo `bittware_xupvvh/software/zcash_fpga.py`

FPGA Memory Map

The FPGA has 2 main methods of sending and receiving data, these are:

1. The AXI4 stream interface, which is used to send and receive commands and can be used with larger amounts of data (detailed in the next section).
2. The AXI4 lite memory map interface, mainly used for configuration, debug, instruction, and data memory. This is done via individual 32 bit writes and reads. The memory space of the FPGA is organized as:

Name	Address range
Top level control and configuration	** Not currently present in version v1.3
Stream control module (only present on AWS builds)	0x0 to 0xFFFF
BLS12-381 coprocessor	0x1000 to 0x4FFF

(each regions memory section is detailed in the architecture section)

Streaming commands

The streaming interface data is streamed from SW to FPGA with a 16 byte header at the very start, and then depending on the command or reply type from FPGA there can be a sub-header and additional inputs / outputs. All values here are little endian and length (len) is specified in bytes. The format of the header is:

```
typedef __packed__ struct {
```

```

uint32_t cmd_type;    // This is the command type (given below) either from SW or
from FPGA
uint32_t len;        // This is the total length in bytes of the packet either from
SW or from FPGA
} fpga_header_t;

```

Commands are capable of being sent back-to-back in the same stream, but the start of a new command must be aligned to an 8 byte boundary.

SW to FPGA

These are the commands the FPGA is capable of receiving from SW.

reset_fpga

cmd_type: 0x00000000

len: 8 (no additional data follows the header)

This command resets the FPGA internal logic logic to its initial state. This should be called when first connecting to the FPGA, or if any errors happen and the FPGA is unresponsive (if this command does not fix the problem you will need to reprogram the FPGA). The FPGA will send a **reset_fpga_rpl** to SW after it has been reset.

get_fpga_status

cmd_type: 0x00000001

len: 8 (no additional data follows the header)

This command asks the FPGA to reply with the current status using a fpga_status_rpl message.

verify_equihash

cmd_type: 0x00000100

len: 8 + 8 + length of block header (CBlockHeader) (1487 for N=200, K=9)

This command takes takes a block header and will verify the equihash solution is correct, according to Zcash protocol doc, and passes the difficulty filter. The FPGA will send a **verify_equihash_rpl** back to SW with the result of the check along with the index from the command so that it can be matched (in the case of multiple concurrent operations).

```

typedef __packed__ struct {
    fpga_header_t hdr;
    uint64_t      index;        // This index is returned with the result
    CBlockHeader  block_header; // Serialized data of block header class from Zcash code
    block.h
} verify_equihash_t;

```

verify_secp256k1_sig

cmd_type: 0x00000101

len: 8 + 8 + 160

This command verifies the signature used in a transparent transaction over the EC **secp256k1**.

Inputs are the hash $H(m)$ of the message m , the signature (comprised of two values - s and r_x), and Q (public key of signer uncompressed). P is the base point of secp256k1 and stored on the FPGA. The FPGA then decodes this command into a series of instructions for the secp256k1 ECDSA core. An index is also given that it returned with the result to track multiple concurrent commands.

```
typedef __packed__ struct {
    fpga_header_t  hdr;
    uint64_t       index;           // This index is returned with the result
    uint256_t      s;              // Signature
    uint256_t      r;              // Signature
    uint256_t      hash;           // Hash of message that was signed to be verified
    uint512_t      Q;              // Signers public key (uncompressed form)
} verify_secp256k1_sig_t;
```

FPGA to SW

These are the replies the FPGA is capable of sending to SW.

reset_fpga_rpl

cmd_type: 0x80000000

len: 8 (no additional data follows the header)

This tells SW that the FPGA has been reset successfully. After this a get_fpga_status message should be sent to the FPGA to confirm it is in a good state.

fpga_status_rpl

Cmd_type: 0x80000001

len: 8 + 36

This reply tells SW the current status of the FPGA, the build information, what commands it is capable of running, and any error flags or extra debug information that might be useful.

```
typedef __packed__ struct {
    fpga_header_t  hdr;
    uint32_t       fpga_version;    // e.g. 0x00_01_00_00 (v 1.0.0, format
    // major.minor.patch)
    uint64_t       fpga_build_date; // String of build date FPGA image was built
    uint64_t       fpga_build_host; // String of machine name FPGA image was built
    uint64_t       fpga_cmd_cap;    // Bitmask of what commands are capable to run on
    // this FPGA build
    uint64_t       fpga_state;      // What the FPGA state is in and any error flags
} fpga_status_rpl_t;
```

fpga_ignore_rpl

Cmd_type: 0x80000002

len: 8 + 8

This reply tells SW that the FPGA received a message it was unable to decode (either did not have the capability or some error in the message, for example incorrect length), and is ignoring it.

```
typedef __packed__ struct {
    fpga_header_t hdr;
    fpga_header_t ignored_command; // This is the command that the FPGA ignored
} fpga_ignore_rpl_t;
```

verify_equihash_rpl

cmd_type: 0x80000100

len: 8 + 8 + 1

This command from FPGA gives the result of a **verify_equihash** command, along with the index and resulting bitmask for any errors found (will be all 0 if it verifies correctly).

```
typedef __packed__ struct {
    fpga_header_t hdr;
    uint64_t      index;
    uint8_t       result_mask; // [0] == DIFFICULTY_FAIL, [1] == XOR_NON_ZERO, [2] ==
BAD_IDX_ORDER, [3] == BAD_ZERO_ORDER;
} verify_equihash_rpl_t;
```

verify_secp256k1_sig_rpl

cmd_type: 0x80000101

len: 8 + 8 + 1

This command replies to SW with the result of the verification check for a secp256k1 signing. We return the result of the verification along with the index. The result passed if none of the result_mask bits are set.

```
typedef __packed__ struct {
    fpga_header_t hdr;
    uint64_t      index;
    uint8_t       result_mask; // [0] == R_OUT_OF_RANGE, [1] == S_OUT_OF_RANGE, [2] ==
X_INFINITY, [3] == SIGNATURE_VERIFICATION_FAILED
} verify_secp256k1_sig_rpl_t;
```

bls12_381_interrupt_rpl

cmd_type: 0x80000200

len: 8 + 4 + N

This command replies to SW when an interrupt instruction is hit by the bls12-381 coprocessor, along with the data that was pointed to by the instruction. The length in the header will for up to the data[N], to know how much data to process in this message you need to parse the data_type.

```
typedef __packed__ struct {
    fpga_header_t hdr;
    uint32_t      index; // Custom value from user in the instruction
    uint8_t       data_type // What type of data (e.g. Affine point, scalar, JB point)
    uint8_t       data[N] // Depending on data in slot this will be from 48 to 576 bytes
    long
} bls12_381_interrupt_rpl_t;
```

FPGA command capability register

This is the bit mask returned from the fpga_status_rpl_t message. If a command is sent to the FPGA for something it has no capability to run, it will reply with a “fpga_ignore_rpl_t”.

Bit	Capability	Note
0	verify_equihash with N= 200, K = 9	Only one of these can be enabled per FPGA build
1	verify_equihash with N= 144, K = 5	
2	verify_secp256k1_sig	Verify a secp256k1 signature
3	BLS12-381 coprocessor enabled	Used to accelerate zk-SNARK

FPGA Architecture

Overview of blocks in the system

These are the blocks in the system, build-time parameters can control which optional blocks are included in the FPGA build (e.g. you might disable those that aren't used so the system fits on a smaller FPGA). Depending on if all blocks are enabled or not, the internet clock speed to FPGA might need to be lower to take into account that the FPGA will have a harder time to close timing constraints.

- Top level board
 - Control block (**required**)
 - Equihash verification engine (optional)
 - Verify pow
 - Find solution (mine)
 - Blake2b for generating XORs
 - SHA256 for difficulty check
 - Hash Map for checking duplicates
 - Order checker of indexes
 - Transparent Signature Verification Engine (secp256k1 ECDSA core) (optional)

- 256b Scalar multiplier mod p / mod n
- 256b Scalar inversion mod p / mod n
- High speed 256b integer multiplier with mod reduction stage of either n or p
- Point add
- Point double
- Point multiply
- Resource arbitrator (to share 256b multiplier core)
- BLS12-381 Coprocessor (zk-SNARK accelerator) (optional)
 - Resource arbitrator sharing
 - 381b integer multiplier mod p
 - 381b integer adder mod p
 - 381b integer subtractor mod p
 - Dual mode F_p / F_{p^2} point operations on bls12-381
 - Point add / double
 - Point multiply
 - Instruction memory
 - Data memory
- Interface module (**required**)
 - UART (For Bittware board)
 - PCIe (For Amazon AWS)

This section talks in more detail about the architecture of each main engine on the FPGA, along with performance results.

Interface module

AWS (Amazon)

The AWS top level has a wrapper `cl_zcash_aws_wrapper.sv` which maps the data coming in over PCIe 512 bits wide to the 64 bits wide expected by the Zcash internal logic. It is also responsible for mapping to the streaming interface. The top level parameter “USE_AXI4” controls if AXI4 or AXI4-lite will be used for the streaming interface.

VHH (Bittware)

This top level has a wrapper to generate the required clocks, and to provide an interface from the USB-UART into the Zcash internal logic.

Equihash Verification Engine

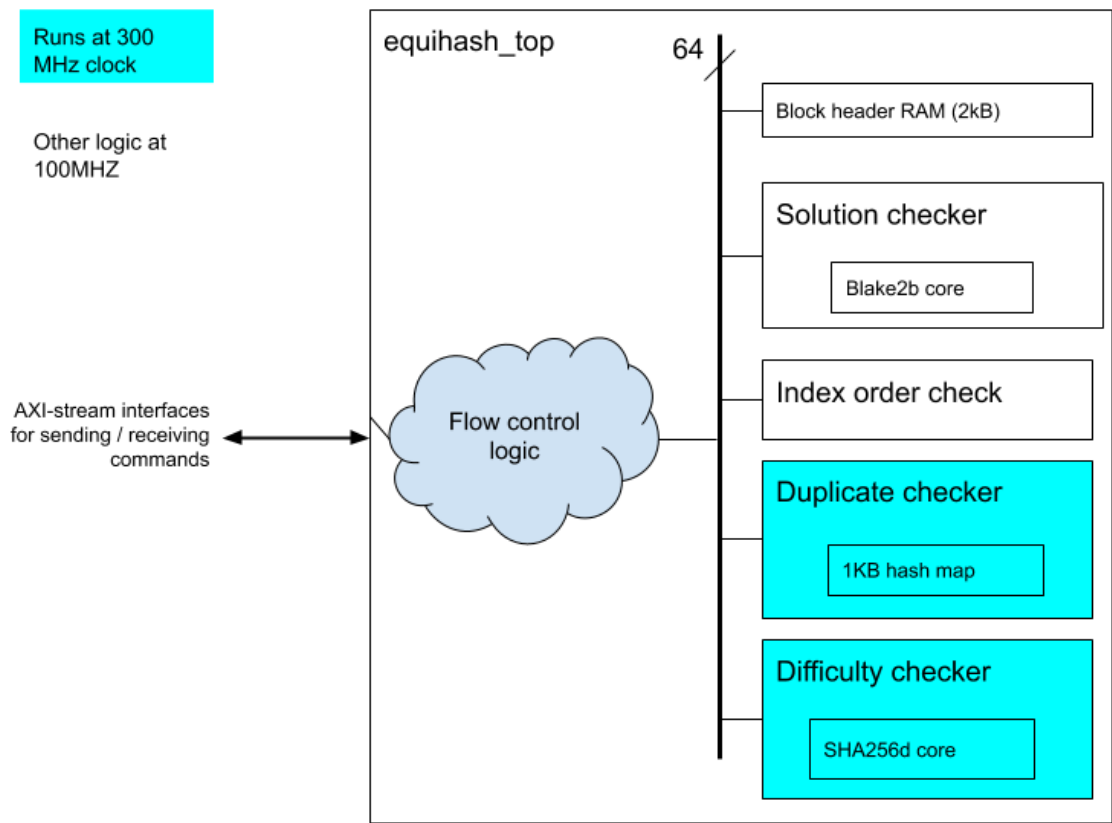
Overview

The equihash engine takes in a block header and then stores the data in global memory, and each sub-block is given the data required for it to check, each which will set a single bit in the resulting block mask. The blake2b block is fully unrolled and running at 200MHz, meaning it takes 64 clock cycles to get a single result, but after that each clock cycle is a new result. This allows the hash of the 512 XOR strings in the equihash solution to be computed at very high throughput. This is more important for parameters ($n=200, k=9$) than the proposed ($n=144, k=5$) as there are less hashes to be performed. The duplicate checker is a hash map and can run at a higher frequency of 300MHz. All the checks run in parallel so the slowest check will determine the

performance, currently the duplicate check and difficulty check. This could be improved by moving both to a higher clock frequency.

The Blake2b core is able to generate a new hash output after an initial delay of $2 + \text{ceil}(\text{input bytes}/128)*24$, so for the solution checker here (140B input, 512 hashes), we achieve 177M hash/s. Maximum performance would be at 5G hash/s.

Block diagram



Performance evaluation

FPGA resources

Percentages reported for the VU37P

LUT	FF	DSP	BRAM
87914 (3%)	54362 (3%)	0	6 (0.2%)

Clock cycles

	FPGA clock cycles	FPGA throughput	CPU cycles	3GHz CPU throughput
Solution check	600 @100MHz			
Index order check	356 @100MHz			
Duplicate check	1443 @300MHz			
Difficulty check	1068 @300MHz			

Equihash solution verification	1068 @300MHz	207K op/s	~2868040	~1K op/s
--------------------------------	--------------	-----------	----------	----------

Here performance on FPGA is 207X faster, likely due to high performance Blake2b core, as well as all checks being done in parallel.

Transparent Signature Verification Engine (secp256k1 ECDSA core)

Overview

This engine handles all the operations for the curve secp256k1. This block at a top level supports point multiplication with a top level state controller, point multiplication, point addition, point doubling, point inversion, integer multiplication, and integer modulo reduction blocks. Blocks are shared via a resource arbitrator.

We optionally can use the endomorphism of secp256k1 to split the k in $X=kQ$ into two smaller half-size k_1 and k_2 , by instantiating a “endomorphism decom block”, which gives close to a 2x improvement in throughput, at the cost of having 2 more multiplication engines.

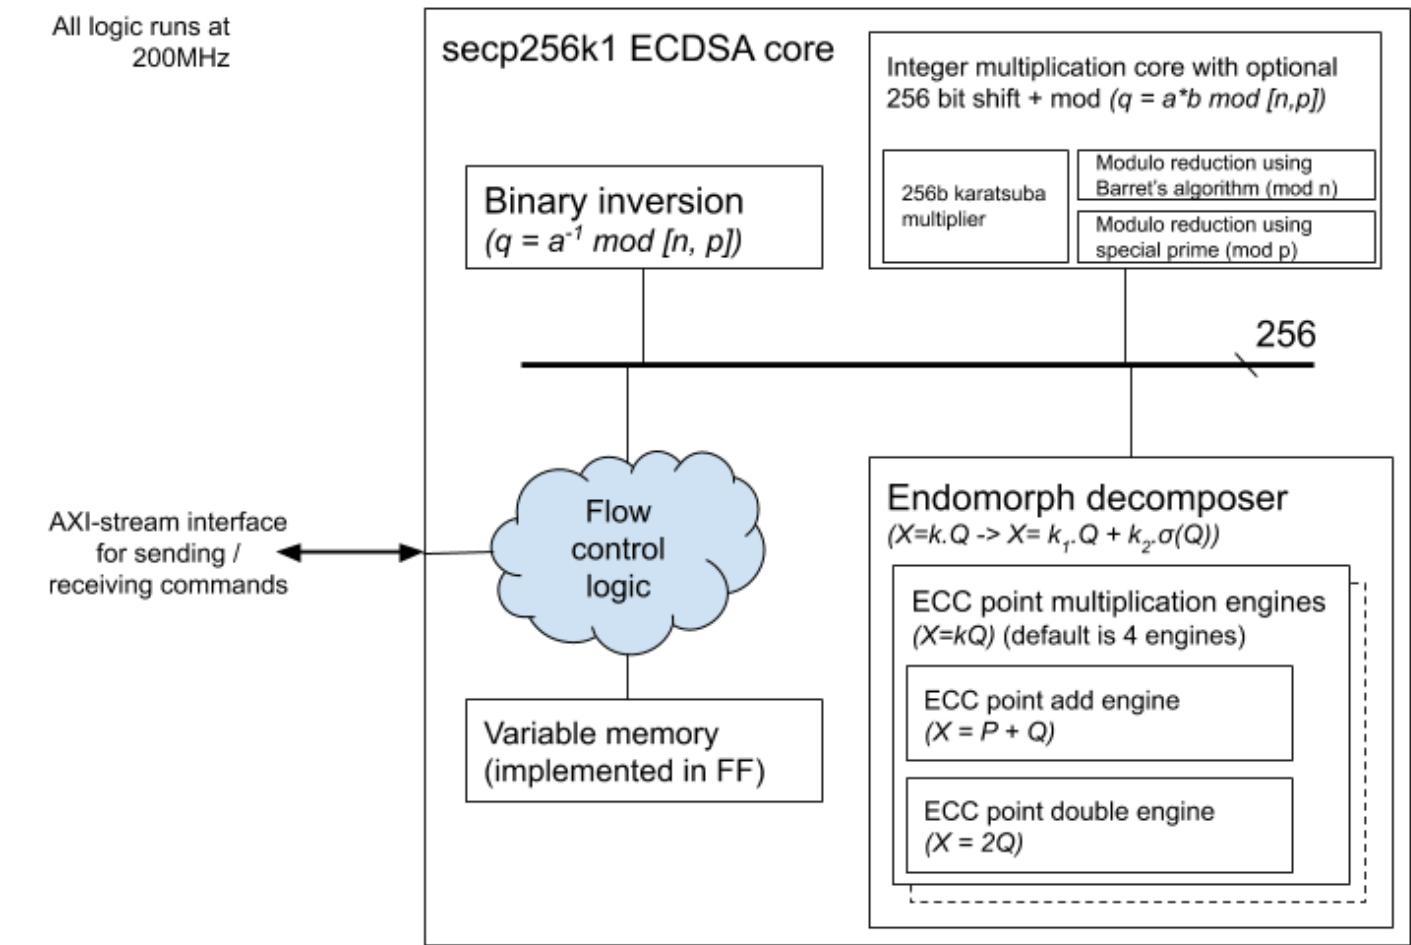
We create two ECC point multiplication modules which run in parallel to calculate $X = u_1P + u_2Q$ required for signature verification. These run in parallel, but due to the pipelined integer multiplication core we have both point multiplication modules share this.

The algorithm used for point multiplication is the double and add method, but we take advantage of the FPGA parallelism and do the double and add at the same time. Since doubling is faster than adding, we start the next double if we have an unfinished add in progress, improving performance. Each point double takes 54 clock cycles and each point addition takes 104 clock cycles.

The integer multiplication core is implemented with the karatsuba algorithm (2 levels) and each level is pipelined over 3 flip-flops (for timing @ 200MHz), so that a result is valid after 6 clock cycles, with a new result every clock cycle. On the output we can optionally bit shift (used for endomorph decomposition), reduce the result mod p (taking 2 clock cycles as it takes advantage of the prime form), or reduce the result mod n which takes longer as it uses barrett's algorithm. Mod n operations are only required at the start and very end so this path does not need to be optimized too much.

Binary inversion uses the gcd algorithm and takes roughly 708 clock cycles, so this is avoided as much as possible by: 1) Converting to jacobian coordinates for point multiplication and 2) the signature result can be checked without converting back to affine coordinates using the same method as in Zcash's git source code.

Block diagram



Performance evaluation

Average performance of the core is shown below for signature verification (which will depend on the number of adds/doubles required). This is also compared to the same function from zcash's git running on a 3GHz processor (measured using average of CPU cycle counts). I did not try to optimize by using non-adjacent form . (NAF) window methods / Shamir's trick, as on the FPGA we run the calculations truly in parallel and might not benefit from these techniques, although this could be a point for future exploration. Improving the equations used for point double and point add would also improve performance. The FPGA was successfully meeting timing at a 200MHz clock. FPGA throughput could be improved by instantiating more cores.

FPGA resources

Percentages reported for the VU37P

	LUT	FF	DSP	BRAM
secp256k1 ECDSA core (without endomorph enabled)	57697 (4.4%)	31751 (1.2%)	144 (1.6%)	2 (0.1%)
secp256k1 ECDSA core (with endomorph enabled)	98792 (7.5%)	61909 (2.1%)	144 (1.6%)	2 (0.1%)

Clock cycles

	FPGA clock cycles	FPGA throughput	CPU cycles	3GHz CPU
--	-------------------	-----------------	------------	----------

				throughput
Point double mod p	54	3.7M op/s		
Point add mod p	104	1.9M op/s		
Inversion mod n	708	282K op/s		
secp256k1 ECDSA core (without endomorph enabled)	20224	9.9K op/s/core	223350	13.4K op/s
secp256k1 ECDSA core (endomorph enabled)	10100	20K op/s/core		

FPGA performance is 1.5X compared to a 3GHz CPU. The FPGA could instantiate multiple ECC engines to run in parallel.

Future Optimizations

Investigating the impact using NAF has on performance would be the next possible optimization.

BLS12-381 Coprocessor (zk-SNARK accelerator)

Overview

This coprocessor is used to accelerate zk-SNARKS as the majority of elliptic arithmetic used during proving and verifying is run on top of the bls12-381 curve.

Unlike previous cores, the coprocessor can be configured by writing to instruction memory rather than accepted hard coded commands. This is to allow more flexibility in how the co-processor is used. SW can either poll registers on the FPGA coprocessor or use interrupt instructions so that the FPGA will send data to SW.

The coprocessor has instruction memory that can be written to, after a reset command the entire memory is initialized to NOOP-WAIT. The coprocessor has a memory bank with addressable data slots each 64 bytes wide per address for variables that can be used with instructions, example sizes for variables are:

- Scalar integer takes 1 slot
- Point in F_p takes 3 in jacobian coordinates (2 in affine)
- Point in F_{p^2} take 6 in jacobian coordinates (4 in affine)
- $F_{p^{12}}$ element takes 12 slots

Each data slot only uses 48 bytes on the FPGA (64 bytes of address space is used in SW to simplify the mapping of memory to slot index). The first 381 bits of a slot store that elements data, the remaining 3 bits are used as a format for the type of element stored (more bits can be added if needed).

0	Scalar
1	F_p element
2	F_{p^2} element
3	$F_{p^{12}}$ element

4	Fp point AF
5	Fp point JB
6	Fp ² point AF
7	Fp ² point JB

Instructions

Instructions are 8 bytes each (1 byte for op-code, and then the rest is used to address variables).

Interrupts are sent by using the SEND-INTERRUPT instruction which can be used to send the result of a calculation to SW. SW will have a method of registering a callback function that would be called when an interrupt is detected, the function will take a pointer to memory that will hold the data sent from FPGA.

Montgomery form is not used in any of the operations (as we can use RAM lookup table technique for the modular reduction).

All point operations can be given inputs in affine or jacobian coordinates, but outputs will be in JB unless otherwise stated. There is not a specific instruction for converting to affine coordinates because you can get the same result by multiplying the point element (Fp or Fp²) by INV-ELEMENT(MUL-ELEMENT(Z, Z)).

Instruction	Description
NOOP_WAIT (0x0)	Coprocessor waits at this command and does nothing (used to stall or after a reset)
COPY_REG(0x1, a, b)	Copy contents of register b = a
JUMP(0x2, a)	Jump instruction pointer to location a
JUMP_IF_EQ(0x4, a, b, c)	Jump instruction pointer to location a if b == c, else go to next instruction (b and c are limited to the lower 64 bits)
JUMP_NONZERO_SUB(0x5, a, b)	If b != 0 then jump to a and b--, otherwise go to next instruction (b is limited to the lower 64 bits)
SEND_INTERRUPT(0x6, a, b)	Send an interrupt to SW along with the data in slot a. Amount of bytes sent will depend on data type stored in slot. 16 bit value of b will be appended to the interrupt message header (see streaming commands for bls12_381_interrupt_rpl_t)
MUL_ELEMENT (0x10, a, b, c)	Do Fp / Fp ² / Fp ¹² field element multiplication, c = a x b
ADD_ELEMENT (0x11, a, b, c)	Do Fp / Fp ² field element addition, c = a + b
SUB_ELEMENT (0x12, a, b, c)	Do Fp / Fp ² field element subtraction, c = a - b
INV_ELEMENT(0x13, a, b)	Calculate the inverse of a Fp / Fp ² field element a and store in b

POINT_MULT(0x20, a, b, c)	Do a F_p / F_p^2 point multiplication using scalar a and F_p / F_p^2 affine point b, and store result jacobian point in c. $c = a \times b$
¹ MILLER_LOOP(0x21, a, b, c)	Do a miller loop of the G1 F_p affine point in a and G2 F_p^2 affine point in b, and store result F_p^{12} field element in c
¹ FINAL_EXP(0x22, a, b)	Do a final exponentiation of the F_p^{12} field element in a and store result F_p^{12} field element in b
ATE_PAIRING(0x23, a, b, c)	Do an optimal ate pairing of the G1 F_p affine point in a and G2 F_p^2 affine point in b, and store result F_p^{12} field element in c

Notes:

¹The purpose of these commands is to allow for faster multi-pairing operations - you can call the MILLER_LOOP instruction on multiple points, and then only a single FINAL_EXP instruction.

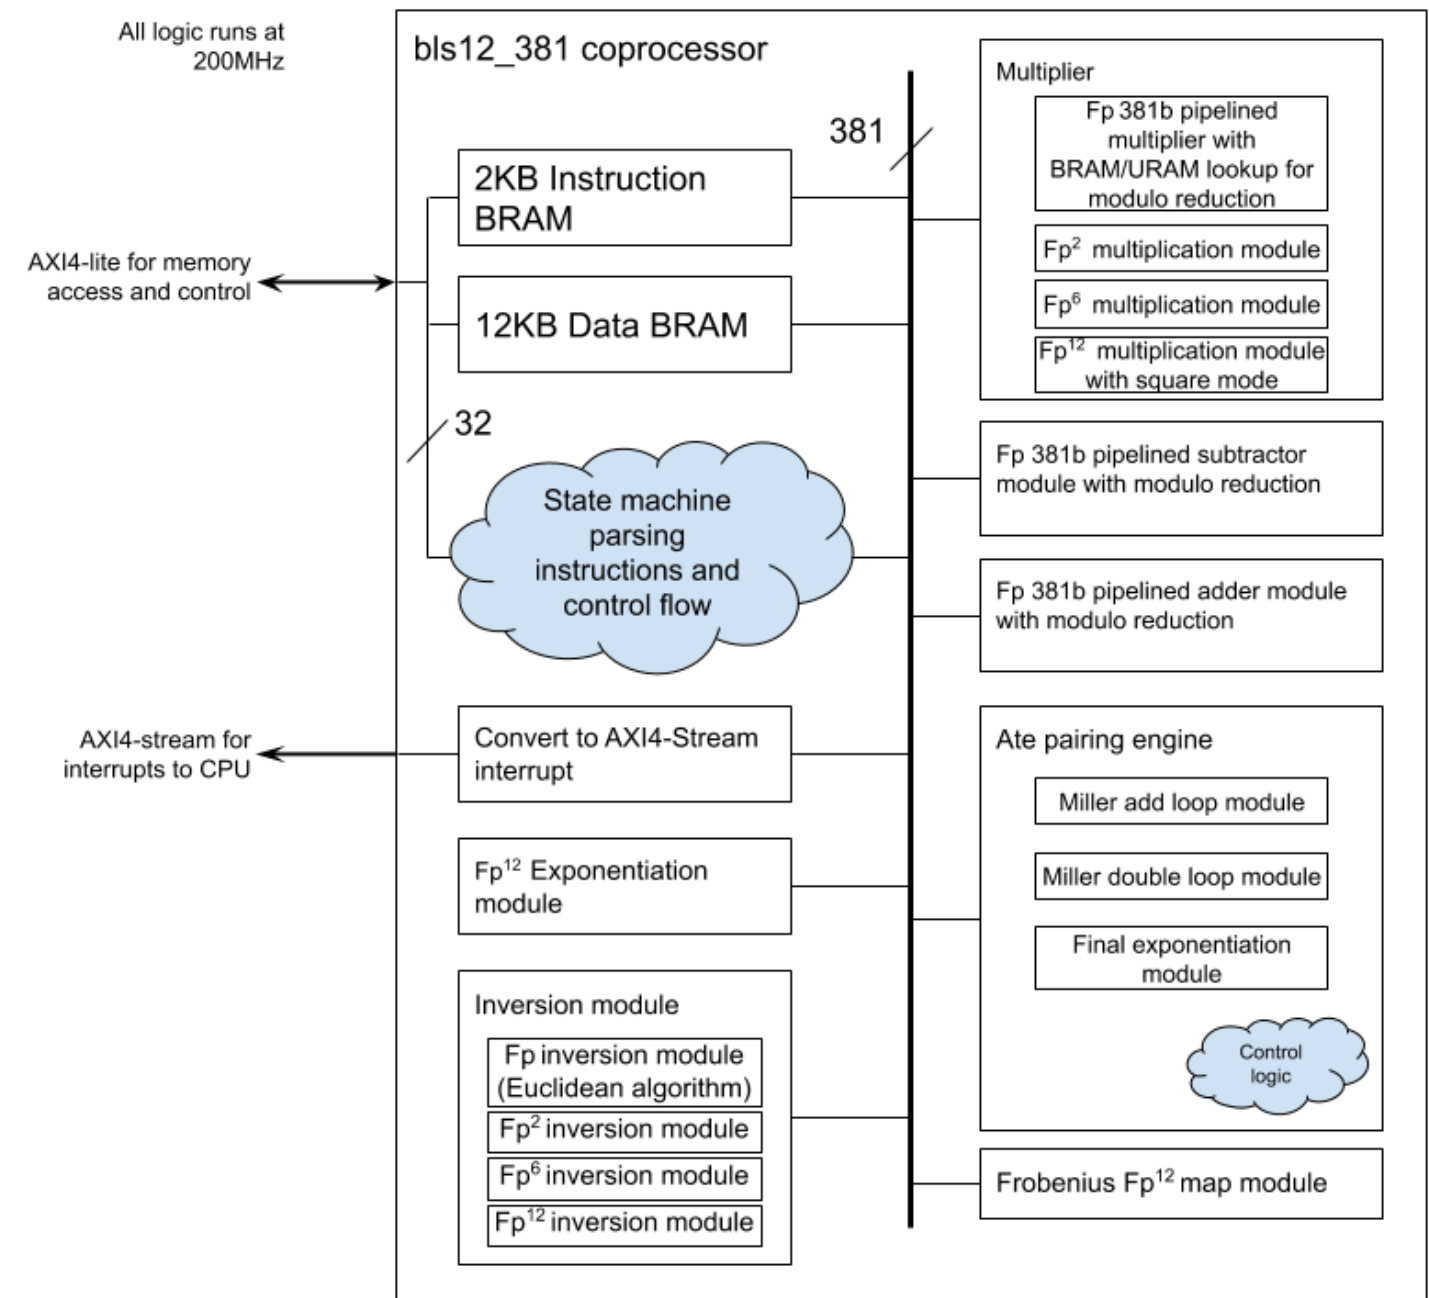
Memory Map

This is the AXI-lite portion of the core that can be used for configuration, as well as writing/reading instruction/data memory.

Register Name	Address	Access
Instruction memory offset / reset control	0x0	Read: returns the memory offset where instruction memory begins Write: A '1' to bit[0] will reset the instruction memory, a '1' to bit[1] will reset the data memory
Data memory offset	0x4	Read only: returns the memory offset where data memory begins
Data memory size	0x8	Read only: returns the power of 2 number of data memory slots (i.e. 8 => 256 slots)
Instruction memory size	0xc	Read only: returns the power of 2 number of instruction memory slots
Current instruction pointer	0x10	Read: returns current instruction memory pointer Write: sets the instruction memory pointer (will wait until current operation finishes)
Last instruction cycle count	0x14	Read only: returns the number of clock cycles the last instruction took to complete
Data for v1.4 multiplier (reduction RAM)	0x18	A write here will load 32 bits onto the reduction RAM data line required for 1.4v multiplier
Control for v1.4 multiplier (reduction RAM)	0x1c	A write to bit[1] will enable a shift of data written to location 0x18.

		A write to bit[0] will enable a write to the reduction RAM address (this should be done after shifting 381bits of data via 0x18 and writes to bit[1]).
--	--	--

Architecture



The coprocessor operates on a shared 381 bit bus with a main state machine with pointers into a data and instruction memory (implemented using Xilinx Ultra RAM on the FPGA). The top level multiplier, adder, and subtractor are all fully pipelined (so a new result each clock) and are resource shared with the entire coprocessor (so inversion block, dual mode point multiplier, pairing engine,... all use this).

The multiplier used is sized so to take advantage of the FPGA DSPs - 27x17 bit wide multipliers. We perform all multiplications in parallel, followed by a tree of carry-save adders. After this we use a RAM lookup technique

for the modular reduction. We then propagate carries, and perform another stage of modular reduction, and final stage of checking if we need to subtract the modulus again.

Point multiplication is supported by either placing a dedicated point multiplication module (same as secp256k1), or by reusing the point double and add modules in the miller loop block which is used for the pairing calculation. Reusing the miller loop logic saves FPGA resources, but results in a slower point multiplication.

Below is a diagram showing the hierarchy of modules and their respective resource usage on the FPGA taken from Vivado. The different main blocks are:

1. The pairing engine, which consists of the Miller loop and final exponentiation.
2. The F_p^{12} inversion module
3. The multiplier - in v1.4 we changed this multiplier to a version that uses RAM lookup, so most of the space is occupied by the carry save adder trees.
4. The F_p^{12} exponentiation module
5. The frobenius map F_p^{12} module

bls12_381_top

627389

bls12_381_pairing_wrapper


453335

accum_mult_mod

151318

bls12_381_pairing

211196



Leaf Cells

52295

(3)

bls12_381_fe12_inv_wrapper

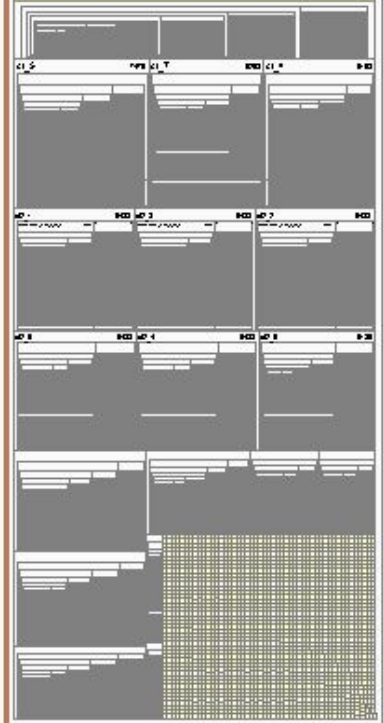
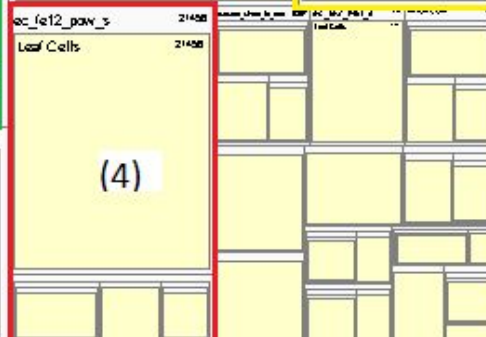
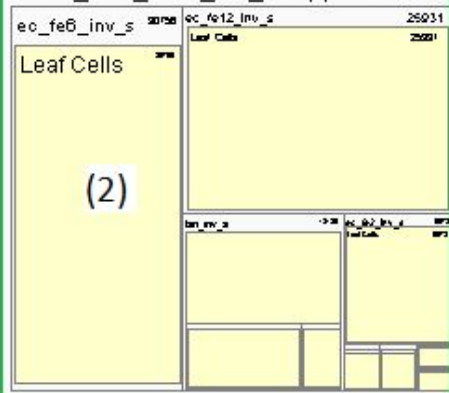
78854

ec_fe6_mul_s

35798

bls12_381_fe12_inv_wrapper

24551



Performance Comparison

FPGA resources

Percentages reported for the VU37P.

LUT	FF	DSP	RAM
327k (25.1%)	226.6k (8.7%)	345 (3.8%)	133 URAM (13.8%), 231 BRAM (11.4%), 14164 LUTRAM (2.3%)

Clock cycles

Here performance was benchmarked vs the Rust bls12_381 crate on a 32GB, 3.7GHz i5-9600K CPU. FPGA is running at 200MHz. Although we have lower throughput for individual operations in F_p , we are able to take advantage of parallel operations inside the higher order F_p^{12} operations, and see a 2.9x speedup in the final ate pairing. A large amount of time is spent on the final exponentiation, which could be a target for optimization - or moving to the weil pairing which does not require the final exponentiation (instead requires two miller loop iterations).

	FPGA clock cycles	FPGA throughput (op/s)	3.7GHz CPU throughput (op/s)
F_p inversion	2685	74.5K	109K
F_p^{12} inversion	3565	56K	60.5K
F_p multiplication + modulo reduction	9	22M	20.8M
F_p^{12} multiplication + modulo reduction	270	740K	228K
F_p point multiplication	49800 (dedicated F_p^2 point mult block)	4016	4926
F_p^2 point multiplication	62064 (dedicated F_p^2 point mult block)	3222	1499
Optimal Ate pairing miller loop stage	38844	5148	1747
Optimal Ate pairing final exponentiation stage	87800	2277	854
Optimal Ate pairing total	126644	1580	553

Future Optimizations

- Investigating the impact of NAF on point multiplication
- Pre-computation for the G2 double / add values used in the miller loop (useful if we are doing multiple pairings)
- Implement sparse multiplication for F_p^{12} for the miller loop
- Modify architecture to use redundant polynomial form, then addition / subtraction / multiplication could be done faster as we don't need to propagate carries.

User Guide

This section goes over example usage of the system.

Running Simulations

Module level simulations

Most modules have a corresponding “_tb.sv” in the tb/ folder, and are self checking so can be added to local copy of Vivado and ran, and will print a message that all tests passed for that module if there are no problems.

A top level simulation that tests all functions and emulates the Bittware VVH top level is here

https://github.com/bsdevlin/zcash-fpga/blob/master/zcash_fpga/src/tb/zcash_fpga_top_tb.sv

The easiest way is to start a new Vivado project, and add all .sv and .xci (ip) files to the project, then run the _tb file you want. A good place to start is

https://github.com/bsdevlin/zcash-fpga/blob/master/zcash_fpga/src/tb/bls12_381_top_tb.sv or

https://github.com/bsdevlin/zcash-fpga/blob/master/zcash_fpga/src/tb/zcash_fpga_top_tb.sv.

When simulating the top modules, they include the multiplier with RAM lookup for modular reduction, so you need to run this script otherwise simulation / building will produce an error:

https://github.com/bsdevlin/zcash-fpga/blob/master/ip_cores/accum_mult_mod/scripts/generate_files.py

It is also recommended defining FASTSIM, otherwise the adder module will consume a lot of time to simulate (https://github.com/bsdevlin/zcash-fpga/blob/master/ip_cores/accum_mult_mod/src/rtl/compressor_tree_3_to_2.sv).

AWS Board level

The simulation test cases for the AWS board are in the repo folder

https://github.com/bsdevlin/zcash-fpga/tree/master/aws/cl_zcash/verif/tests and can be run by:

1. `cd /home/centos/aws-fpga/hdk/cl/developer_designs/cl_zcash/verif/scripts`
2. `make all`

This will compile the test cases and then run them, they are all self checking, so if not ERRORS are printed and the simulation finishes then there are no problems. If something unexpected happens you can run xsim and look at the waveforms. At the moment there is just one test (test_zcash.sv) that will test all of the top level block functions..

Usage with a local FPGA board

If the board is local, it can be configured over USB-UART (note this is very low bandwidth and just mainly used for proof of concept / testing).

Commands can be called from the python script: `bittware_xupvvh/software/zcash_fpga.py`

Usage on AWS

AWS runs over PCIe and has a higher bandwidth, but due to timing a slower clock is used (as there is more glue logic on the FPGA).

At the time of writing this these were the versions used in the AWS toolchain:

Developer Kit Version (HDK)	Tool Version Supported (Vivado)	Compatible FPGA developer AMI Version
1.4.8-1.4.X	2018.3	v1.6.0 (Xilinx SDx 2018.3)

Building the FPGA image

If you make changes to the code or want to build a new image, you can follow the steps below. If you do not want to do this, you can skip to the next section “Loading FPGA image” and use one of the pre-built images listed in “Existing AFIs”. When building the image it is recommended to change the parameters here to only enable the blocks required -

https://github.com/bsdevlin/zcash-fpga/blob/3a8c799a742061760d9c1deaaaebd72a60792ca9/zcash_fpga/src/rtl/top/zcash_fpga_pkg.sv#L32 , as enabling everything will make the build take longer and might not meet timing.

1. Start an AWS instance and load it with the FPGA Developer AMI (<https://aws.amazon.com/marketplace/pp/B06VVYBLZZ>)
 - a. This should be a f1 instance (e.g. f1.2xlarge) so you have access to an FPGA
 - b. If you just want to build the FPGA image you can use a cheaper instance like r5.xlarge (just need at least 32GB RAM)
2. Clone the zcash git repo
3. Clone the aws-fpga repo
 - a. `git clone https://github.com/aws/aws-fpga.git`
4. Copy the folder `zcash-fpga/aws/cl_zcash` to the AWS folder
 - a. `cp -r /home/centos/zcash-fpga/aws/cl_zcash /home/centos/aws-fpga/hdk/cl/developer_designs/`
5. Copy the folder `/home/centos/aws-fpga/hdk/cl/examples/common`
 - a. `cp -r /home/centos/aws-fpga/hdk/cl/examples/common /home/centos/aws-fpga/hdk/cl/developer_designs/`
6. Run the `hdk_source.sh` script to setup the AWS environment
 - a. `cd /home/centos/aws-fpga; source hdk_setup.sh`
 - b. Note: If you get an error with Vivado not being present, it might be due to locale issue, try:
 - i. `export LC_ALL="en_US.UTF-8"`
7. Set the variables for Zcash scripts:
 - a. `export CL_DIR=/home/centos/aws-fpga/hdk/cl/developer_designs/cl_zcash; export ZCASH_DIR=/home/centos/zcash-fpga/`
8. Generate the FPGA IP files
 - a. `cd /home/centos/aws-fpga/hdk/cl/developer_designs/cl_zcash/ip/; ./run_cl_sde_ip_flow`
9. Start building the FPGA image
 - a. `cd /home/centos/aws-fpga/hdk/cl/developer_designs/cl_zcash/build/scripts; ./aws_build_dcp_from_cl.sh -clock_recipe_a A0 -clock_recipe_b B1`
 - b. Note: AWS clock recipes are here: https://github.com/aws/aws-fpga/blob/master/hdk/docs/clock_recipes.csv , a higher performance version of the core can use “`-clock_recipe_a A1 -clock_recipe_b B0`”, a slower version (but

easier to build and meet timing) could use “*-clock_recipe_a A2 -clock_recipe_b B1 -strategy BASIC*”

- i. Note: this will not work with an ILA debug core since the clock speed (15MHz) is too slow compared to JTAG frequency
- c. You can check progress by looking at
“/home/centos/aws-fpga/hdk/cl/developer_designs/cl_zcash/build/scripts/last_log”

The build should run and will take several hours, depending on the instance type / clock recipe. If there are no problems, the output will be in

/home/centos/aws-fpga/hdk/cl/developer_designs/cl_zcash/build/checkpoints/to_aws/*.tar and needs to be uploaded to an Amazon S3 bucket. The bucket used in this project is “zcash-fpga-west”. From here you can follow the standard flow detailed on the AWS FPGA github:

<https://github.com/aws/aws-fpga/blob/master/hdk/README.md#step3> .

After this you should have an agfi-ID that can be used to program the FPGA.

Loading FPGA image

To load an FPGA image you need it's agfi-ID, either from the previous step or from the table in the following section “Existing AFIs”.

Run this commands to load the FPGA:

1. `sudo fpga-load-local-image -S 0 -I -F agfi-ID`

Note: You can check for errors / metrics by running the command “`sudo fpga-describe-local-image -S 0 --metrics`”. If you see all 0's then there is no problem, but if you see some timeouts like this:

```
ocl-slave-timeout-addr=0x2001
```

```
ocl-slave-timeout-count=4
```

You should reload the FPGA image (step 1 above). There is a known issue with AWS where the first load will sometimes show this problem, but reloading FPGA fixes it.

Rust interface

A rust interface has been developed to allow the Zcash client to utilize the FPGA acceleration.

Startup test program

A simple program is in /home/centos/aws-fpga/hdk/cl/developer_designs/cl_zcash/software/runtime/(test_zcash)

1. Run the sdk_source.sh script to setup the software AWS environment
 - a. `cd /home/centos/aws-fpga; source sdk_setup.sh`

Run the make file and then the test program (using sudo), and check there are no errors. The test program connects to the FPGA and programs the bls12-381 coprocessor to do a multi-pairing and then send the result

back to SW, and then tests the secp256k1 ECDSA verification (will first check these are enabled in the build). Expected output would be this (taken from v1.3.0 build):

[illegible]

```

slot 4, pt: 3,
data:0x1409cebef9ef393aa00f2ac64673675521e8fc8fddaf90976e607e62a740ac59c3dddf95a6de4fba
15beb30c43d4e3f8
slot 5, pt: 3,
data:0x1692a61ce5f4d7a093b2c46aa4bca6c4a66cf873d405ebc9c35d8aa639763720177b23beffaf522d
5e41d3c5310ea333
slot 6, pt: 3,
data:0x081abd33a78d31eb8d4c1bb3baab0529bb7baf1103d848b4cead1a8e0aa7a7b260fbe79c67dbe41c
a4d65ba8a54a72b6
slot 7, pt: 3,
data:0x0900410bb2751d0a6af0fe175dcf9d864ecaac463c6218745b543f9e06289922434ee446030923a3
e4c4473b4e3b1914
slot 8, pt: 3,
data:0x113286dee21c9c63a458898beb35914dc8daaac453441e7114b21af7b5f47d559879d477cf2a9cbd
5b40c86becd07128
slot 9, pt: 3,
data:0x06d8046c6b3424c4cd2d72ce98d279f2290a28a87e8664cb0040580d0c485f34df45267f8c215dcb
cd862787ab555c7e
slot 10, pt: 3,
data:0x0f6b8b52b2b5d0661cbf232820a257b8c5594309c01c2a45e64c6a7142301e4fb36e6e16b5a85bd2
e437599d103c3ace
slot 11, pt: 3,
data:0x017f1c95cf79b22b459599ea57e613e00cb75e35de1f837814a93b443c54241015ac9761f8fb20a4
4512ff5cfc04ac7f
slot 12, pt: 3,
data:0x079ab7b345eb23c944c957a36a6b74c37537163d4cbf73bad9751de1dd9c68ef72cb21447e259880
f72a871c3eda1b0c
INFO: All tests passed!

```

FPGA debug

Debug instructions can be found here:

https://github.com/aws/aws-fpga/blob/master/hdk/docs/Virtual_JTAG_XVC.md

There is a parameter in cl_zcash.sv, USE_ILA = "NO" which can be changed to "YES" to enable a build with the debug logic. You can change the connections as needed.

Latest AFIs

These are the latest bug-free (no known bugs) public AFIs that exist and can be used on an AWS F1 instance.

agfi-ID	afi-ID	Notes
agfi-0d0aeee105030594a	afi-081b60a4044e3db15	v1.3, Contains secp256k1 sig core and bls12-381 coprocessor

Old (unused) AFIs

These are listed here for tracking purposes but not intended to be used as they are mostly debug / have bugs.

agfi-ID	afi-ID	Notes
agfi-0528daff45454ed7c	afi-09056704c94b5280b	v1.0.0 First test version used for testing AWS flow, will not work with test program.
agfi-05561b352d56b5f57	afi-0c8109482d730073c	v1.0.1 Test version
agfi-0fa84678db6b2752f	afi-07ec21206df23e398	v1.1.0, Has all modules enabled but on a slow clock recipe for testing. BLS12_381 core has Fp and Fp2 fpoin instructions
agfi-019c2736fd0141219	afi-0b891a8fc9644f1a0	v1.1.0_150, only has BLS coprocessor enabled but running at 125MHz, uses AXI4 as PCIe interface
agfi-05468e41c302eb331	afi-06a4b56d6e4bfd896	v1.1.1, contains all cores @ 125MHz, uses AXI-lite as PCIe interface
agfi-0fce4c1ad9e0c6c43	afi-0da67f631a2573656	v1.1.2 contains all cores @ 125MHz
agfi-0c4a39d7638bc6010	afi-0bcef9f0c08bee7c1	v1.1.2 contains all cores @ 15MHz
agfi-0abc260b651d87d41	afi-0075820f5d00bd799	v1.1.3 Bug fixes to BLS12_381 core, 125MHz
agfi-07ae22f20d6e90559	afi-0e49dd7ef17fda51a	v1.1.4, bug fix for multiple back to back interrupts, 125MHz
agfi-0db37e1358c1d885f	afi-0907df570f7dc7b2b	Debug version of v1.1.4 above (15MHz)
agfi-06d033b207d8f65c5	afi-07177e176d04aa84b	v1.1.5, debug version 125MHz, contains bug fix for inverter, fp^12 logic

Conclusions

We were able to realize the main goals of this project:

- Accelerate blockchain verification
 - We developed an equihash verification engine that is able to take an input block header, verify the equihash solution is correct, matches the required difficulty, ordering requirement, and index uniqueness. With the current equihash parameters of $N=200$, $K=9$ we were able to achieve 207x speedup compared to the current Zcash SW client. The parallelism of the FPGA was able to be exploited fully here, as we can calculate many Blake2b hashes in parallel required for the solution verification, as well as doing the other checks at the same time. The engine takes parameters for the equihash values of N and K so can be adapted to other settings (e.g. $N=144$, $K=5$).
 - For verifying transparent transactions on the Zcash blockchain we developed a secp256k1 ECDSA core. This core is able to verify the signature used on a transparent transaction, which is the same as a transaction in Bitcoin. We were able to achieve 1.5x speedup when compared to the current Zcash SW client. The main reason the speedup is less when compared to the equihash verification is the lack of parallelism that we can exploit for a single signature verification.
- Accelerate zk-SNARKs
 - We developed a bls12-381 coprocessor which is able to perform curve operations that are required for zk-SNARKs. The coprocessor was designed with a simple instruction set so that it can be programmed from SW and is flexible in the flow of operations it can perform. All EC operations have also been implemented as SystemVerilog software models in the bls12_381_pkg.sv file, allowing for testbench verification and ease of implementation. We achieved 3x speedup compared to a 3.7GHz processor.
- Develop an open source FPGA code base for benefiting the wider community
 - All the code developed has been made open source and released under the GNU General Public License v3.0. It has all been developed in SystemVerilog (the most modern hardware description language), and uses parameters where possible so that the code can easily be reused and can be of benefit to the greater community (other zk-SNARK projects, other crypto coins, research projects...). Even after this project finishes it is expected that the code will continue to be worked on and improved in the open source community.

We decided not to implement FFT acceleration on the FPGA as after talking with Zcash engineers it was decided the benefit to zk-SNARKs would be minimal when compared to accelerating EC pairing operations. Although this could be a future project as FPGAs have historically been used to calculate FFTs and there exists a large amount of reference IP and code.

When developing code for operations on the bls12-381 curve it was evident the main consumer of FPGA clock cycles was the Fp multiplication operation followed by the modular operation over a prime where we could not use any special tricks to modular reduce in HW. Time taken to optimize this area of the code, or different prime selection would greatly speedup the overall performance of operations.

We accelerated the optimal ate pairing, but we found the time to calculate the final exponentiation was more than 2x the time for a miller loop on the FPGA (could be due to not enough optimization in $F_p / F_{p^{12}}$ multiplication), but this points that possibly a different pairing could be fast - such as the weil pairing where we have no final exponentiation stage, but require multiple miller loops.

The bls12-381 coprocessor uses a $F_p \rightarrow F_{p^2} \rightarrow F_{p^6} \rightarrow F_{p^{12}}$ tower, as the Zcash bls12-381 Rust create code was used as a reference when developing the FPGA code, but it would be worth investigating if different towerings would be faster on FPGA - it's proposed $F_p \rightarrow F_{p^2} \rightarrow \mathbf{F_{p^4}} \rightarrow F_{p^{12}}$ tower might be faster.

Appendix

Example decoding Zcash block #346

Hash 0x000000eff179fb1e47b7aa8667ad4d8e1ef3dbb0d79144030482bf93b5e6339f

Hex dump of block (CBlock):

0	04	00	00	00	13	d6	d1	a4	10	51	42	19	f7	2f	f3	a0	df	d5	c3	8b	62	1c	c2	c6	68	78	4d	2f	d6	fd	10	8f	
20	48	00	00	00	30	16	31	55	23	12	34	9d	d5	3b	6b	9e	23	1d	f8	bc	b8	c2	d3	32	64	cc	02	f5	cd	d9	a9	69	
40	fb	93	80	50	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
60	00	00	00	00	3f	85	13	58	bf	c3	03	1e	1b	b2	b5	50	a4	01	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
80	00	00	00	00	00	00	00	00	00	00	00	03	fd	40	05	00	9d	fa	04	89	e1	18	99	dc	5e	50	5d	91	24	57	44	49	
a0	28	12	b3	f6	0e	31	04	e2	1e	98	b2	7d	80	38	c4	41	82	ca	de	1a	e2	ec	dc	e2	77	10	4f	9f	a6	5d	d4	b6	
c0	a1	b6	ab	44	66	24	ef	6a	0a	c2	a8	5e	2e	a3	32	19	7c	d3	cd	51	b6	e8	a3	31	d4	04	d4	68	bc	ed	6b	e6	
e0	19	e4	8f	0b	8f	c4	3b	f9	dd	44	b2	f1	05	b4	7c	b8	e7	e5	eb	21	96	bd	12	89	0e	df	36	07	41	1e	55	81	
100	b9	14	c2	91	b2	a7	1f	27	19	79	7c	bc	49	13	42	34	62	bd	11	fb	d7	b8	00	31	85	01	31	4f	2b	4e	15	1a	
120	87	f6	40	16	9f	d2	19	ef	51	bd	9c	19	94	38	7c	69	88	bf	68	77	7d	69	e3	06	2f	dc	61	0e	4a	43	99	04	
140	b7	d1	f6	26	78	fb	e8	a9	2f	f1	2a	38	0c	b0	5b	29	33	4b	37	7c	c5	30	11	e5	db	a4	80	01	30	56	b5	a5	
160	71	1b	10	e6	35	1e	5d	72	f5	4b	93	63	3b	0e	5d	4c	0b	12	ff	9b	d6	20	31	84	5f	47	fb	90	23	af	db	3c	
180	15	bd	4a	51	ab	b9	9a	d8	0d	4c	ef	21	b5	c9	da	e9	a3	a5	61	a8	97	74	c2	ff	4e	3d	89	92	20	94	37	b7	
1a0	63	f8	9d	22	61	6b	01	15	62	12	f7	40	47	ba	a3	43	6a	a0	5e	bf	3a	25	d5	b9	df	f7	d9	d0	b7	e0	ba	43	
1c0	83	9c	1d	00	b2	3c	01	d1	e9	a8	42	95	06	8f	65	20	fb	53	59	d2	f7	c9	b2	60	44	ab	0f	0f	de	de	8a	02	
1e0	45	62	8d	43	4a	64	bd	96	8d	93	8f	22	6f	a9	75	32	ec	b7	a0	af	27	06	0a	aa	7f	97	3a	2d	b2	95	83	20	
200	35	de	d2	92	4b	08	bc	6a	4a	06	f8	b1	d4	db	b8	55	c1	f0	37	01	db	ba	a7	55	52	93	c4	3a	86	9c	23	3e	
220	3f	2c	c7	50	14	bd	2c	ef	23	aa	ad	e5	1b	3e	d9	08	fc	7b	1a	03	c7	a2	d8	71	8d	16	37	97	28	52	af	95	
240	64	23	21	c5	57	7d	14	80	14	fd	68	e0	a0	96	87	03	c7	7a	d1	8b	7a	ad	29	99	a6	78	d6	0f	63	04	8f	33	
260	30	ff	d3	1c	bb	75	3c	c6	66	c8	35	1f	35	cf	ac	76	46	93	0b	b7	1c	17	8f	86	05	ff	7e	6f	a1	94	71	c8	
280	e1	09	cc	59	13	61	62	07	8b	17	e5	e5	e7	4f	db	49	01	c4	6a	17	2e	25	15	6d	bd	35	43	87	39	f3	a4	da	
2a0	ec	96	ea	dc	fc	78	a4	77	9a	dd	07	26	70	f6	5f	6d	d1	0c	74	96	5c	f3	8b	f2	f2	d6	85	42	b4	54	99	d4	
2c0	58	f5	2d	c9	25	63	35	9c	87	47	48	90	f6	dd	47	61	d8	24	76	6e	f6	4f	07	fd	5b	5c	38	12	ed	9c	b4	4d	
2e0	85	69	47	e0	c2	b2	02	f4	b9	fa	7d	ce	c3	da	05	03	53	6d	a5	1d	65	99	92	19	72	25	96	2c	b6	63	2c	c1	
300	c9	ff	91	35	e2	20	a3	d9	33	ff	8d	fa	2b	24	61	12	93	ad	ae	45	99	76	1b	2e	0e	32	2a	36	7c	a3	ea	f5	
320	44	33	da	78	95	27	53	6d	d5	6a	26	c7	f9	5f	b7	01	cf	9e	2f	00	52	68	11	70	fa	95	50	ad	69	bd	5e	15	
340	f6	9c	81	5f	1b	c7	f7	79	fa	18	30	47	dd	86	f4	61	b1	a3	e3	3b	97	ec	3d	59	b3	17	c4	8d	36	de	ba	7d	
360	8d	fc	d6	e3	71	a8	d9	32	1e	7e	d7	79	c0	a4	44	66	44	16	15	2c	ad	f5	e1	17	64	ba	f0	5f	11	79	cb	8f	
380	fa	4c	42	0a	d3	5f	b5	d8	f4	39	73	b9	c7	33	da	e1	e5	55	1a	57	00	14	fc	03	4f	08	ff	76	4c	64	b5	e1	
3a0	c9	7d	75	d7	a5	40	49	7b	01	66	9f	d3	ec	25	55	69	f9	64	4a	2f	5f	7d	82	36	1a	08	d7	dd	46	35	8f	79	
3c0	47	3e	6b	5d	62	c0	37	66	5e	7b	c0	94	69	30	84	ff	7b	1f	76	60	dd	77	e8	03	fa	95	75	e0	5d	3d	43	fb	
3e0	e3	d0	74	0c	11	ee	51	eb	f1	af	9b	47	08	98	f7	1f	75	4a	7a	d9	bf	5e	f1	7a	f2	14	4c	dc	95	4e	4f	69	
400	e8	13	b0	0f	5f	e9	4e	93	1d	b2	b3	37	cd	10	44	c3	e7	50	e0	9b	68	b2	18	e1	41	5e	25	54	4c	b9	52	83	
420	65	96	0b	e4	bf	02	62	c3	5e	6d	f3	0f	35	85	5e	5e	2f	09	63	8a	14	61	20	1b	0d	53	1e	53	42	96	ba	19	
440	12	dc	73	d0	5d	a3	de	37	9e	f4	b2	c2	40	3b	2b	41	e6	57	d6	45	37	11	03	09	ad	e0	1b	40	78	fe	d6	c2	
460	da	cc	31	05	3e	9d	28	ff	cf	a4	13	db	62	8a	68	2e	95	1f	88	23	63	9a	a7	d1	1b	9d	79	60	b1	ac	35	04	
480	4f	bb	c8	3e	d4	5f	2e	a6	9c	b4	4b	1c	a5	f9	89	fa	9e	ba	fc	23	2e	44	45	0a	c8	55	44	9b	aa	53	d6	f2	
4a0	39	f3	a6	5a	1f	59	d3	3f	06	1e	c4	14	35	db	63	48	cf	ef	df	d4	0b	4c	42	20	6f	16	63	5a	82	b6	25	9b	
4c0	52	d9	ec	0f	0f	9f	0a	fb	85	6a	2c	e0	6f	fa	23	29	9e	b4	0f	05	db	50	74	02	83	28	6b	a9	ba	71	b4	20	
4e0	bd	47	c5	18	c3	af	7c	eb	15	a9	05	3f	26	d7	de	a7	89	31	11	9b	1c	58	64	ae	a4	96	3a	55	6b	06	50	84	
500	36	6b	8a	cc	2d	36	7a	d2	2c	7f	5a	ec	d2	2d	1c	d1	c3	57	2a	2e	52	bf	26	cb	46	00	e0	d7	05	85	ae	38	
520	a7	12	94	78	78	d4	38	07	8c	59	a0	1d	5f	34	f3	6c	08	c1	87	97	5e	98	b4	a7	9b	b3	93	37	12	16	72	d6	
540	ef	cf	39	6b	f8	33	12	d8	9b	51	b4	15	d7	71	3c	f3	5b	19	ea	e7	ae	71	4b	50	93	7e	ee	11	a1	9e	38	58	
560	a8	98	0a	4c	1b	52	33	24	b9	9f	08	e0	a2	d1	a2	2b	93	47	e2	43	fb	ad	a5	38	1a	fe	0a	09	40	fc	ca	b0	
580	ca	34	52	c2	6f	15	b3	82	f3	67	bb	23	89	7e	fe	fd	19	30	f8	db	53	9a	ec	d9	32	ea	c6	46	32	c1	d2	4a	
5a0	61	42	de	11	6d	49	3d	7d	1c	33	14	b0	37	56	cc	07	0d	53	3b	cc	62	6d	2f	bf	38	a2	59	d4	33	f2	cb	5b	
5c0	52	d1	65	66	f6	a9	f5	79	bb	87	18	bf	7b	d5	db	01	01	00	00	01	00	00	00	00	00	00	00	00	00	00	00	00	00
5e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	ff	ff	ff	ff	04	02	5a	01	00	ff	ff
600	ff	ff	02	20	fa	07	01	00	00	00	00	23	21	02	7a	46	eb	51	35	88	b0	1b	37	ea	24	30	3f	4b	62	8a	fd	12	
620	cc	20	df	78	9f	ed	e0	92	1e	43	ca	d3	e8	75	ac	88	fe	41	00	00	00	00	00	17	a9	14	7d	46	a7	30	d3	1f	
640	97	b1	93	0d	33	68	a9	67	c3	09	bd	4d	13	6a	87	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00

Header:

Version:

04 00 00 00

Previous block hash:

13 d6 d1 a4 10 51 42 19
f7 2f f3 a0 df d5 c3 8b
62 1c c2 c6 68 78 4d 2f
d6 fd 10 8f 48 00 00 00

Merkle Root hash:

30 16 31 55 23 12 34 9d
d5 3b 6b 9e 23 1d f8 bc
b8 c2 d3 32 64 cc 02 f5
cd d9 a9 69 fb 93 80 50

Final sapling root hash:

00 00 00 00 00 00 00 00
00 00 00 00 00 00 00 00
00 00 00 00 00 00 00 00
00 00 00 00 00 00 00 00

Time:

3f 85 13 58

Bits (Difficulty):

bf c3 03 1e

Nonce:

1b b2 b5 50 a4 01 00 00
00 00 00 00 00 00 00 00
00 00 00 00 00 00 00 00
00 00 00 00 00 00 00 03

Equihash solution (the 0xfd4005 here is used to decode the length of the array of bytes, 0xfd means the size is stored as a 2 byte integer 0x4005 == 1344 bytes):

fd 40 05 .. 7b d5 db(1344 bytes until address 0x5ce)

Transactions:

Transaction input array size (one transaction):

01

Version (only 4 bytes here as is not overwinter):

01 00 00 00

Input to transaction array size (one input):

01

OutPoint:

00 00 00 00 00 00 00 00
00 00 00 00 00 00 00 00
00 00 00 00 00 00 00 00
00 00 00 00 00 00 00 00
ff ff ff ff

Script (first byte is length, 4 bytes long):

04
02 5a 01 00

Sequence:

ff ff ff ff

Transaction output array size (two transactions):

02

1st transaction output amount (17300000, 0.173 ZEC):

20 fa 07 01 00 00 00 00

1st transaction output script (first byte is length, 35 bytes long):

23
21 02 7a 46 eb 51 35 88
b0 1b 37 ea 24 30 3f 4b
62 8a fd 12 cc 20 df 78
9f ed e0 92 1e 43 ca d3
e8 75 ac

2nd transaction output amount (4325000, 0.04325 ZEC):

88 fe 41 00 00 00 00 00

2nd transaction output script (first byte is length, 23 bytes long):

17
a9 14 7d 46 a7 30 d3 1f
97 b1 93 0d 33 68 a9 67
c3 09 bd 4d 13 6a 87

Locktime:

00 00 00 00